



# PIMP for Feature Importance

A practical guide to understanding what drives your data science models

# Attendance: "Permutation"



# What is PIMP?

Permutation Importance in Machine Learning

A way to determine which features have the MOST impact on your model's predictions

---

**How:** Measure performance loss when you randomly shuffle a feature's values





# Why Use PIMP?



## Model-Agnostic

Works with any machine learning model—from linear regression to deep neural networks



## Captures Complexity

Identifies feature interactions and non-linear relationships that simpler methods miss



## Stakeholder-Friendly

Easy to interpret and explain to non-technical audiences

# Common Use Cases



## Feature Selection

Reduce dimensionality while preserving model performance



## Model Debugging

Validate assumptions and catch unexpected dependencies



## Business Insights

Understand which factors truly drive outcomes



## Data Quality

Identify problematic features or data issues

# How PIMP Works: The Process



## Train & Baseline

Train your model and calculate baseline performance on validation data



## Shuffle Feature

For each feature, randomly shuffle its values while keeping all other features unchanged



## Measure Impact

Calculate model performance with the shuffled feature



## Calculate Importance

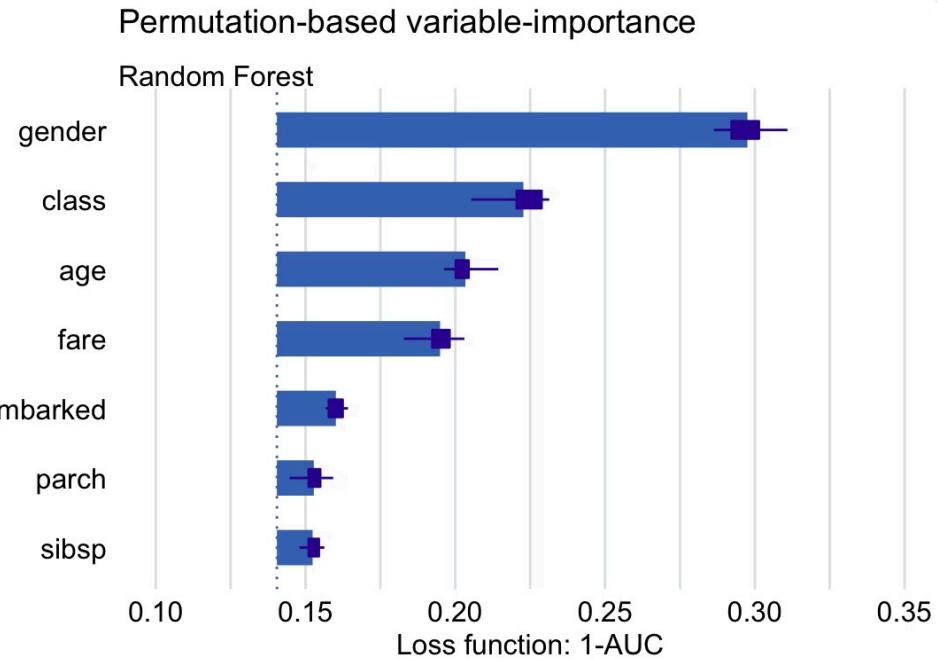
Feature importance = Baseline performance – Shuffled performance



## Repeat & Stabilize

Repeat multiple times (3-50 iterations) to get stable, reliable estimates

# Interpreting PIMP Scores



## High Positive Importance

Model relies heavily on this feature. Shuffling it significantly hurts performance.

**Action:** Keep and protect this feature

## Near-Zero Importance

Feature provides little value to predictions.

**Action:** Consider removing to reduce complexity

## Negative Importance

Feature may be causing overfitting or adding noise.

**Action:** Strong candidate for removal

# PIMP vs Other Methods

## Coefficient-Based Importance

**Limitation:** Only works for linear models

**PIMP advantage:** Works for any model and captures non-linear effects

## Tree-Based Importance

**Limitation:** Biased toward high-cardinality features

**PIMP advantage:** Reflects actual predictive performance

## SHAP Values

**Limitation:** Computationally expensive at scale

**PIMP advantage:** Faster for large datasets; SHAP excels at instance-level insights

- **Best practice:** Use PIMP alongside other methods for comprehensive model understanding



# Best Practices

## ✓ Do

- Use Validation Data

Always use held-out validation data, never training data

- Repeat Permutations

Run 10-50 iterations for stable estimates

- Consider Cost

Account for computational expense on large datasets

## □ Don't

- Confuse with Causation

Importance ≠ causality; correlation can mislead

- Ignore Domain Knowledge

Statistical results must align with business understanding

- Rely on PIMP Alone

Never use PIMP scores as sole decision factor for high-stakes choices



# Common Pitfalls



## Unrealistic Permutations

Shuffling may create feature combinations that don't exist in reality

**Solution:** Validate with domain expertise



## Computational Cost

Can be expensive for large datasets or complex models

**Solution:** Use strategic sampling

"The best defense against pitfalls is combining PIMP with domain knowledge and triangulating with other explainability methods"

# Key Takeaways

## ○ Performance-Based Measurement

PIMP measures feature importance by evaluating performance drop when features are shuffled

## ○ Universal Applicability

Model-agnostic approach captures real predictive value across any ML algorithm

## ○ Essential Toolkit

Powerful tool for feature selection, model interpretation, and debugging

## ○ **Simple concept, powerful insights**

# Google Colab



# Thanks NDL!

1

Next Meeting: 10/26

Topic: SHAP